

Day 2 - 8th February



Multi-Cloud – BiBiGrid

09.00 - 09.30 (EET)

[ELIXIR Germany](https://elixir-europe.org/elixir-germany)

Xaver Stiensmeier, Jan Krüger



Motivation

- **Resource availability**
 - scaling beyond one cloud site
 - making special cloud resources (databases, hardware) available
- **Security** - keep confidential data on a trusted cloud

Why BiBiGrid?


- **Easy to use** - simple tool for cluster setup; multi-clouds become more accessible for research

Our Usecase

- OpenStack on physically separated clouds
- [BiBiGrid](#) setup on an authorized machine
- Direct network access to master node
- distributed Nextflow workflows

For this project we used two locations of [de.NBI Cloud](#)

- [Giessen](#)
- [Bielefeld](#)



BiBiGrid

[BiBiGrid](#) is a cloud cluster creation and management tool for OpenStack (others to come)

- Easily **shareable cluster configurations**
- Uses **clouds.yaml** (provided by many infrastructures like OpenStack, AWS, ...)
- “**On demand**” node scheduling bounded by the initial BiBiGrid configuration
- Has a [Hands-on tutorial](#) (currently without multi-cloud)



openstack
CLOUD SOFTWARE

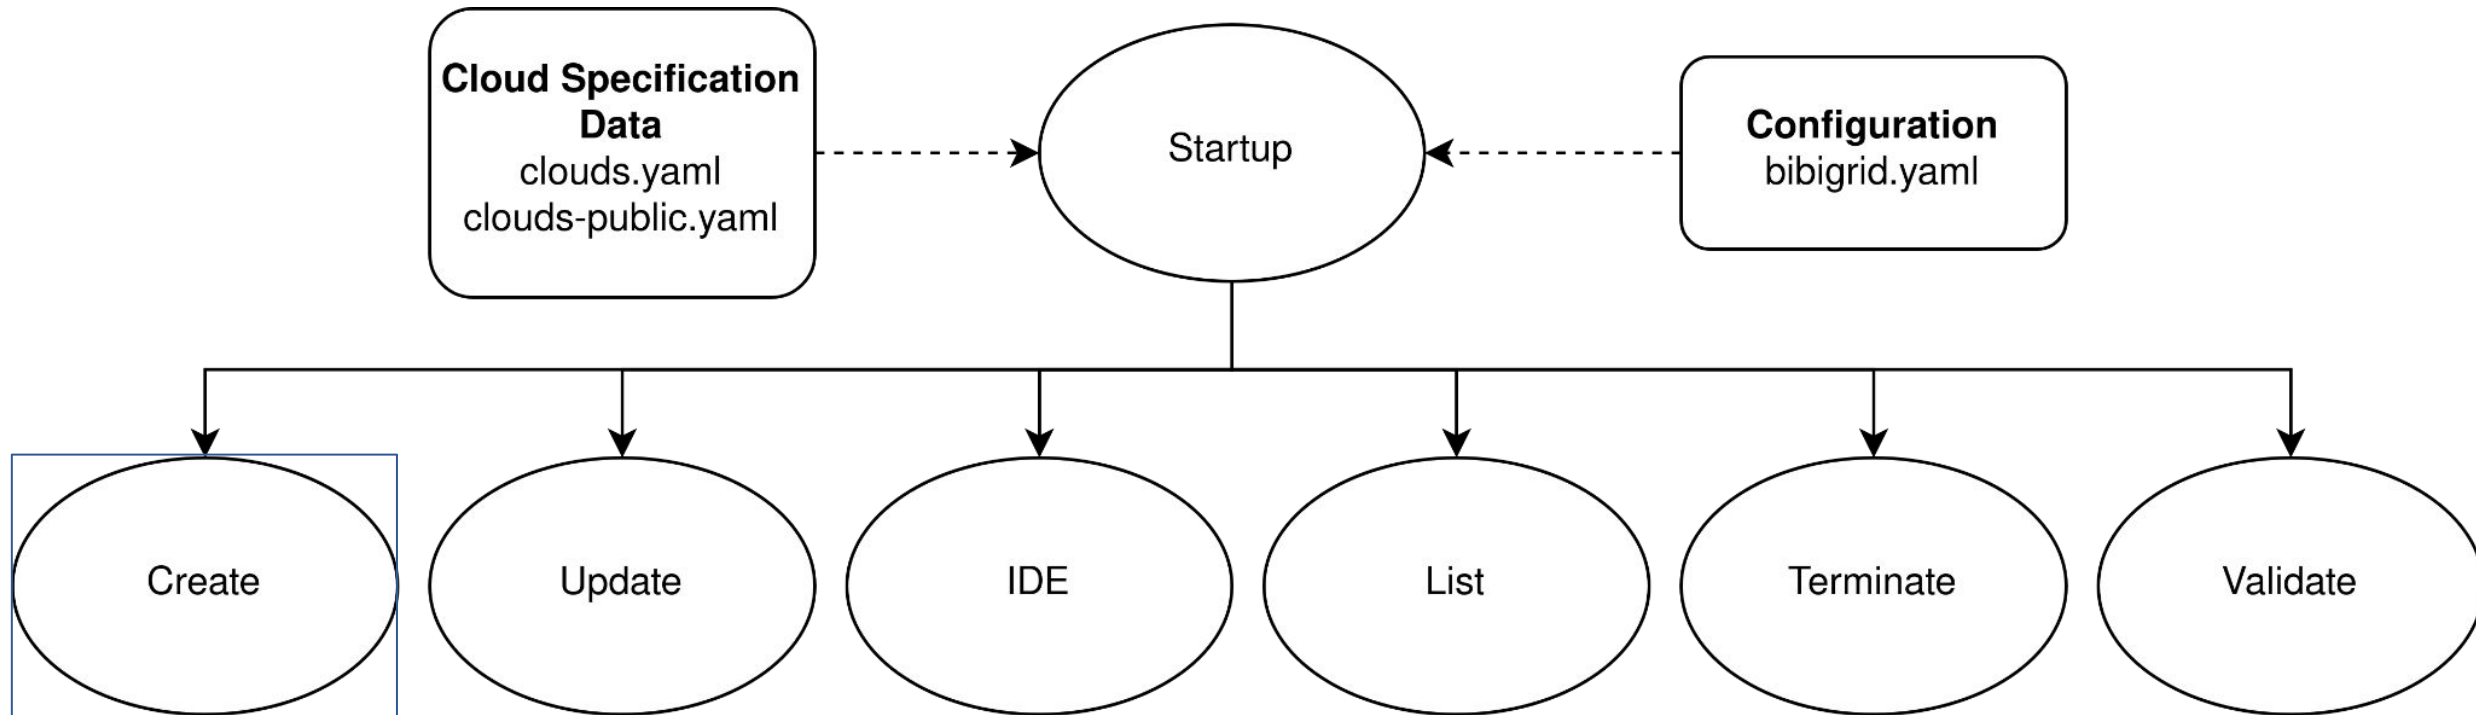


ANSIBLE

slurm
workload manager



BiBiGrid



BiBiGrid Provides

[Ansible](#) – simple server **configuration**

[Slurm](#) – cluster management and **job scheduling** system

[NFS](#) (Network File System) – **shared file access**

[Docker](#) – **run** applications **independent** of the device they run on using **containers**

[Theia Web IDE](#) – easy, intuitive and abstract **server web access**

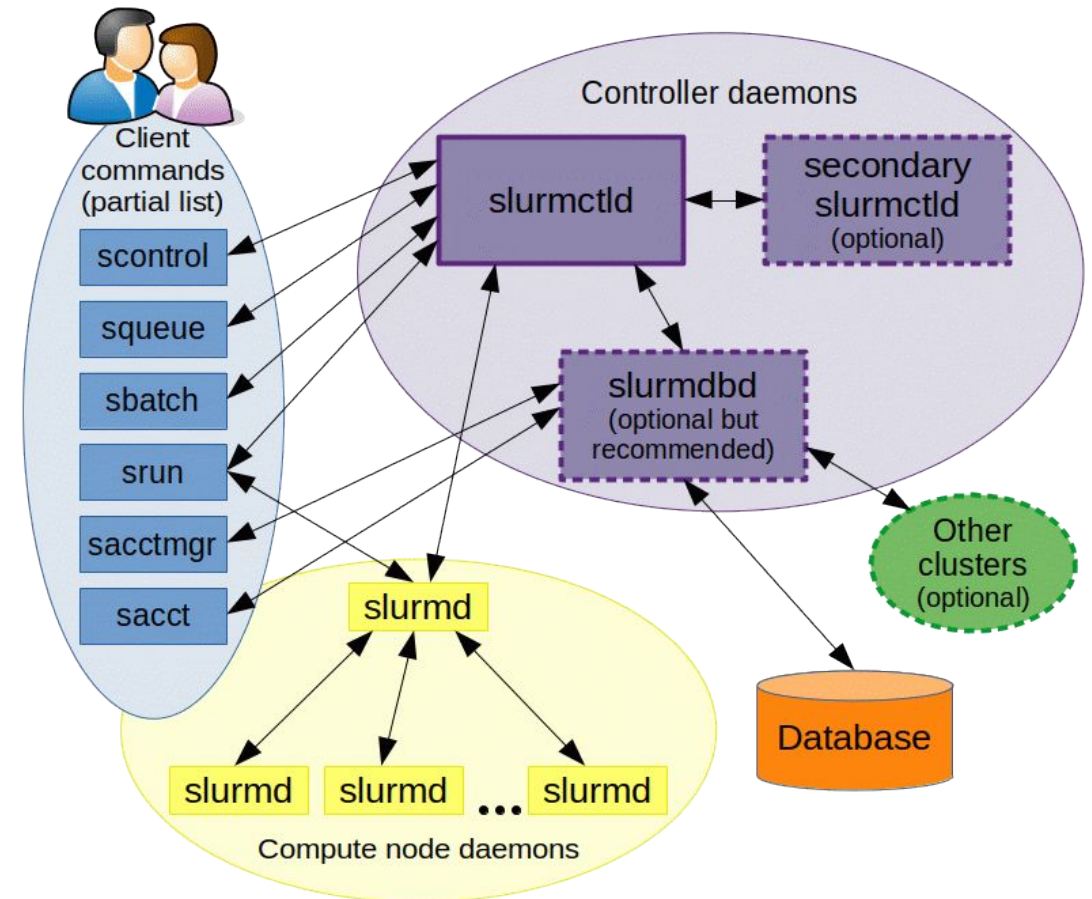
[Zabbix](#) – **monitoring** solution for servers and more

Slurm

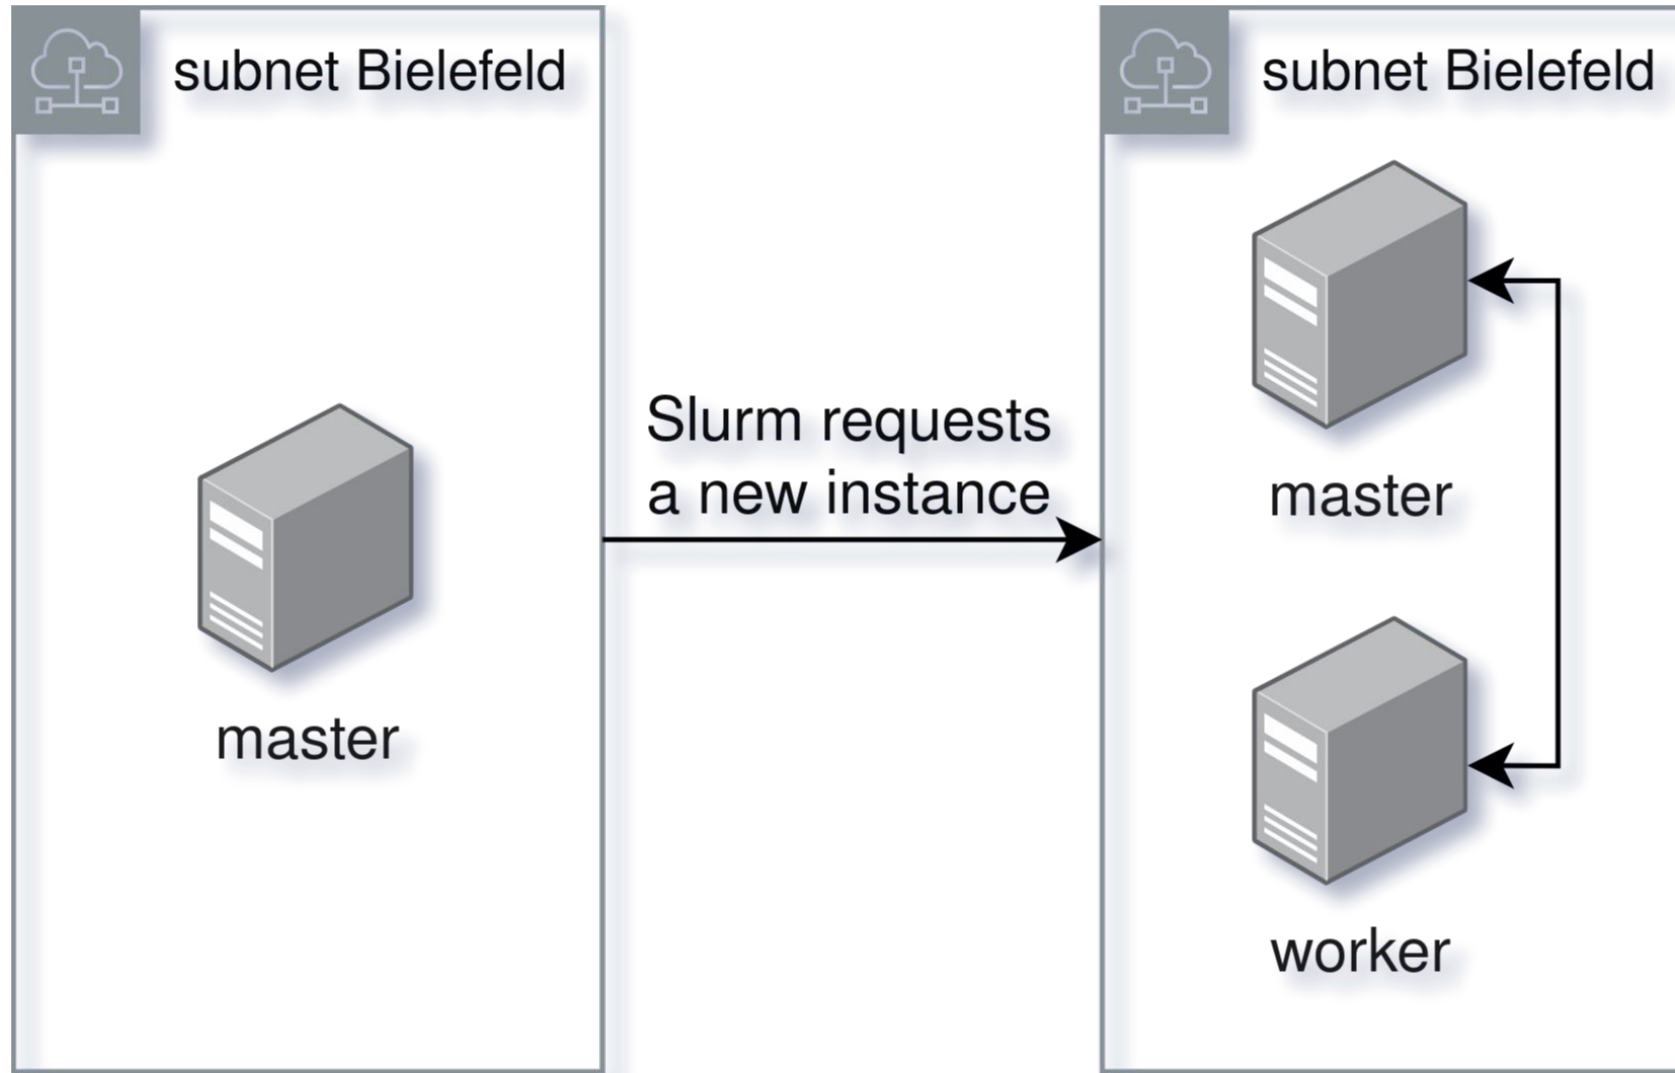
Slurmctld: Management Daemon on master node

Slurmd: Compute Node Daemons running on workers

srun: Submits a job for execution passing along resource and job information



BiBiGrid: "On Demand"



Multi-Cloud – BiBiGrid

- Cloud locations must be able to **communicate**
- Slurm must be able to request server instances for jobs “on demand” **on all clouds**
- Slurm must be able to schedule jobs to **specific clouds**
- Communication between clouds must be **secure**

Multi-Cloud – BiBiGrid

- Cloud locations must be able to **communicate**
- Slurm must be able to request server instances for jobs “on demand” **on all clouds**
- Slurm must be able to schedule jobs to **specific clouds**
- Communication between clouds must be **secure**

- Go easy on resources
- Working on multiple clouds should not complicate use

Slurm - Scheduling Jobs to specific clouds

```
srun job.py
```

→ Command `job.py` executed wherever possible

```
srun --partition Bielefeld job.py
```

→ `job.py` runs in **Bielefeld**

```
srun --partition Giessen job.py
```

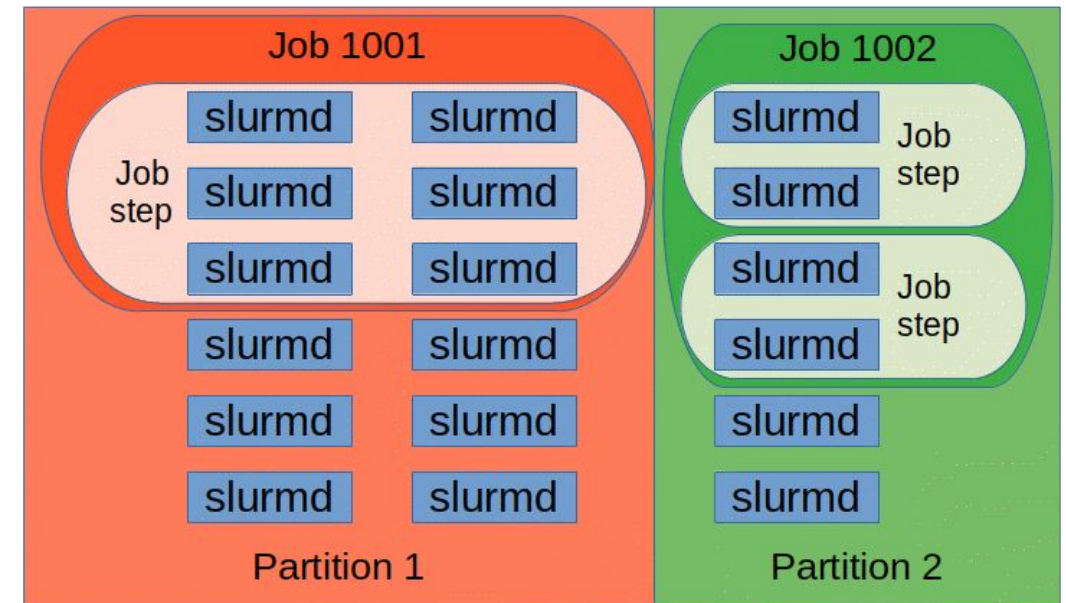
→ Command `job.py` executed in **Giessen**

```
srun --mem 2048G job.py
```

→ runs on a node with at least 2 TB memory

```
srun --cpus-per-task 64 job.py
```


→ tasks require 64 cpus



Bielefeld

Giessen

Idea: Virtual Private Network





```
root@peerA:~/a — Konsole
peerA # wg
interface: wg0
  public key: 0jZc5ZDVoejDyBPX8ZZ50wazzGcBy/nlfRAqKwaMokU=
  private key: MJF54tkjb3wrXZIo2G0b0euYu3Cz/ua5l5Y6bHDBmmY=
  listening port: 51820

peer: JkcUEgA9oqPHClu2l/j04dpBlxkipTG7skRoTB1FnH0=
  endpoint: 192.168.1.2:51820
  allowed ips: 10.0.0.2/32
  latest handshake: 6 seconds ago
  bandwidth: 377 B received, 520 B sent
peerA #
```

```
root@peerB:~/b — Konsole
peerB # wg
interface: wg0
  public key: JkcUEgA9oqPHClu2l/j04dpBlxkipTG7skRoTB1FnH0=
  private key: sA6ljUB+0+nAmeAcU5iJ8xZrtWkiia//mkui9Q021E=
  listening port: 51820

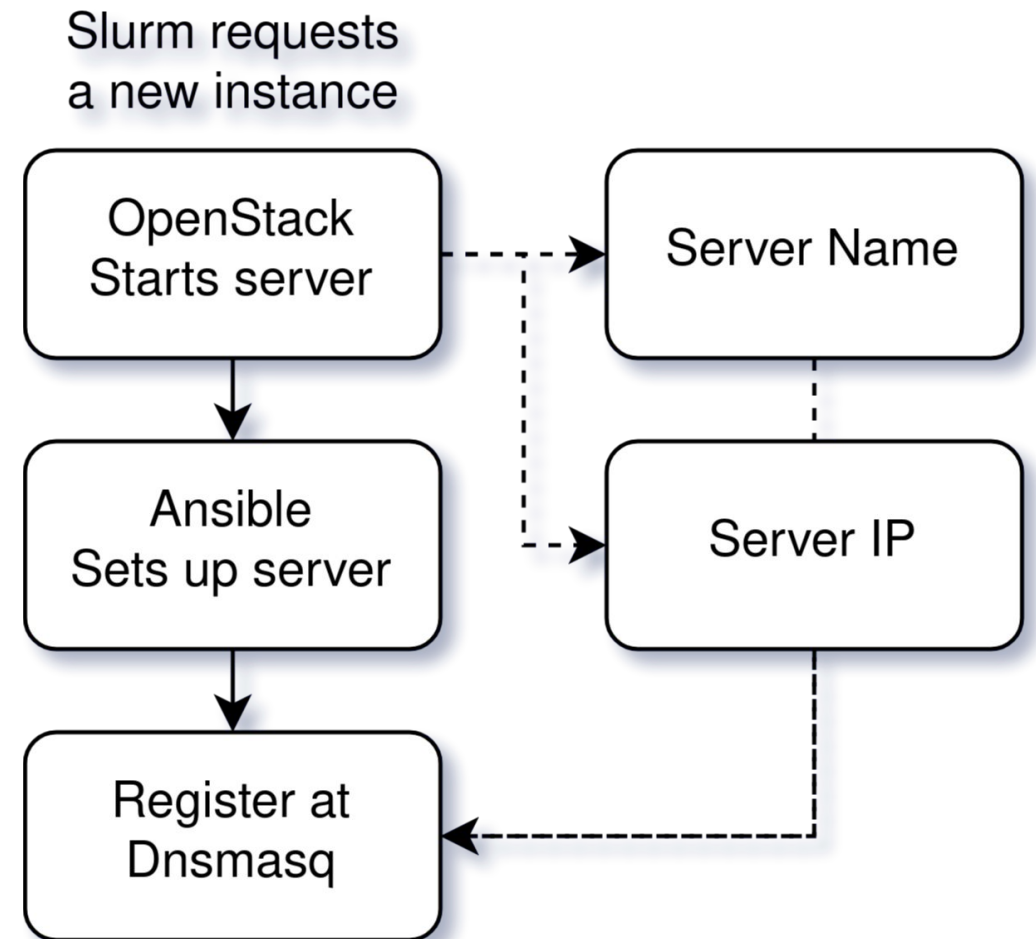
peer: 0jZc5ZDVoejDyBPX8ZZ50wazzGcBy/nlfRAqKwaMokU=
  endpoint: 192.168.1.1:51820
  allowed ips: 10.0.0.1/32
  latest handshake: 5 seconds ago
  bandwidth: 433 B received, 464 B sent
peerB #
```

Dnsmasq - DNS subsystem for a local DNS

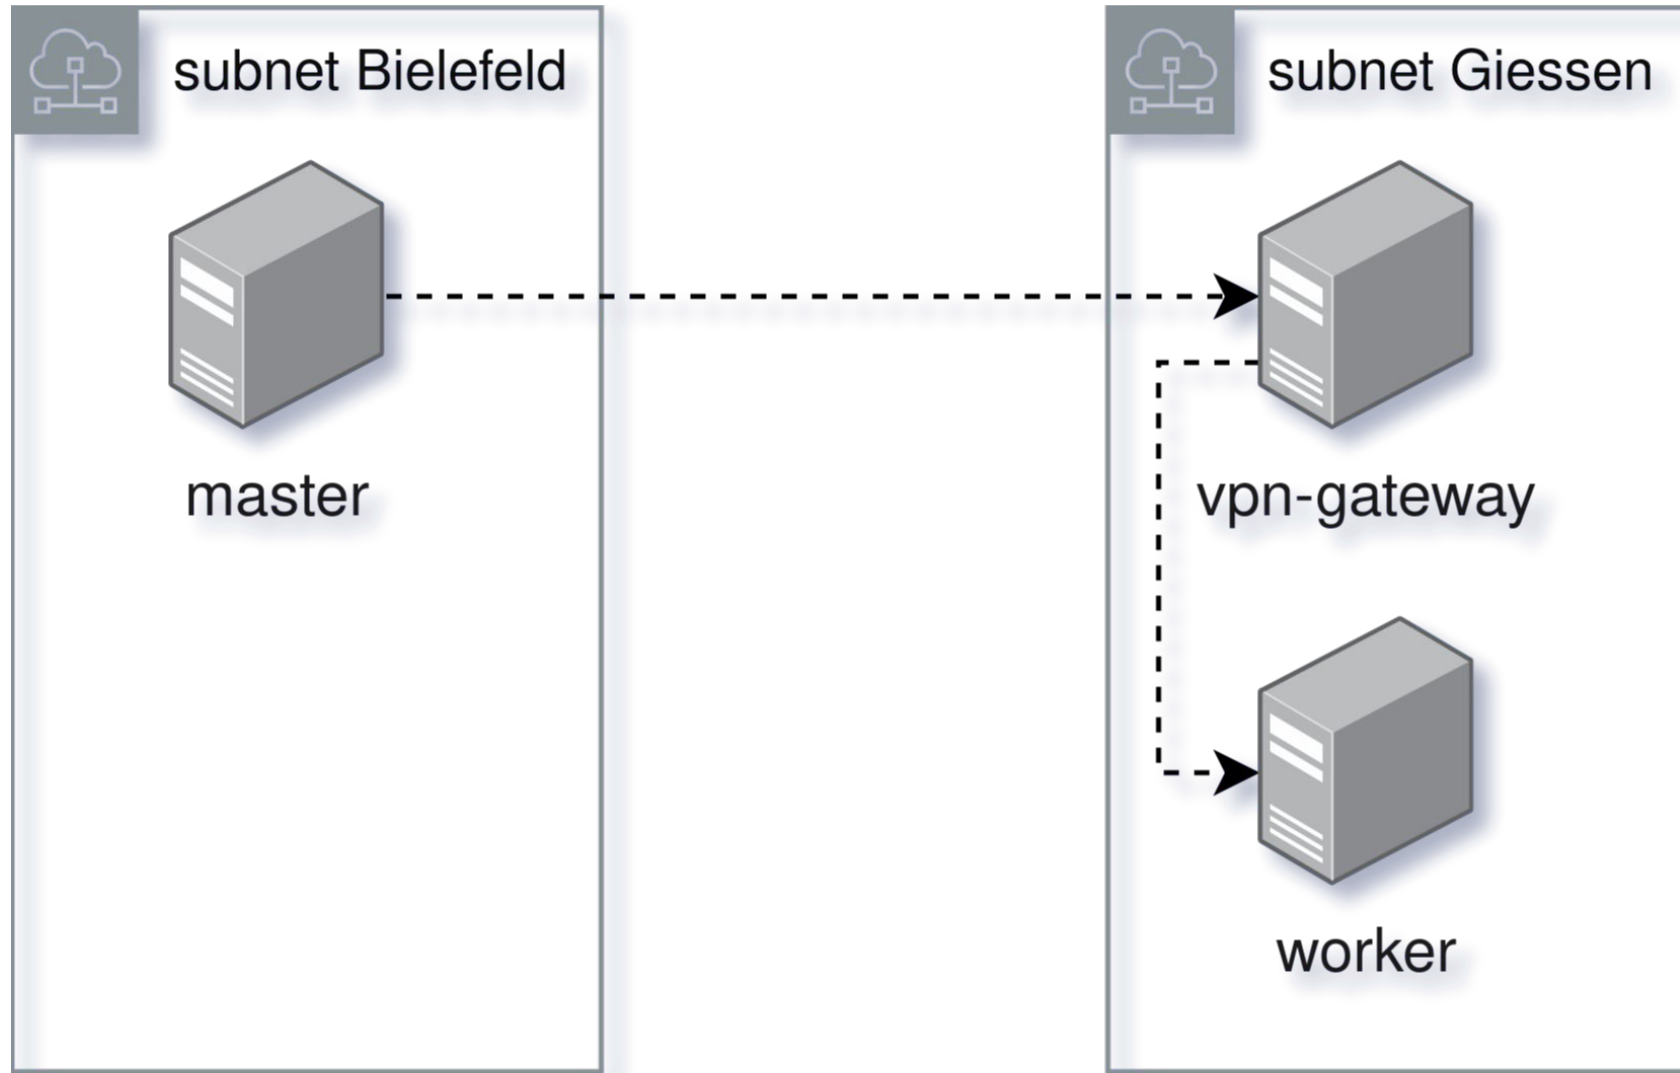
BiBiGrid needs to **provide name resolution** itself – as OpenStack's name resolution only works on its own location.

Name resolution must adapt **"on the fly"** as BiBiGrid provides "on demand" node scheduling using Slurm's [Elastic Computing](#)

=> Configure [Dnsmasq](#) whenever a new instance is requested



Idea: Virtual Private Network



Routing

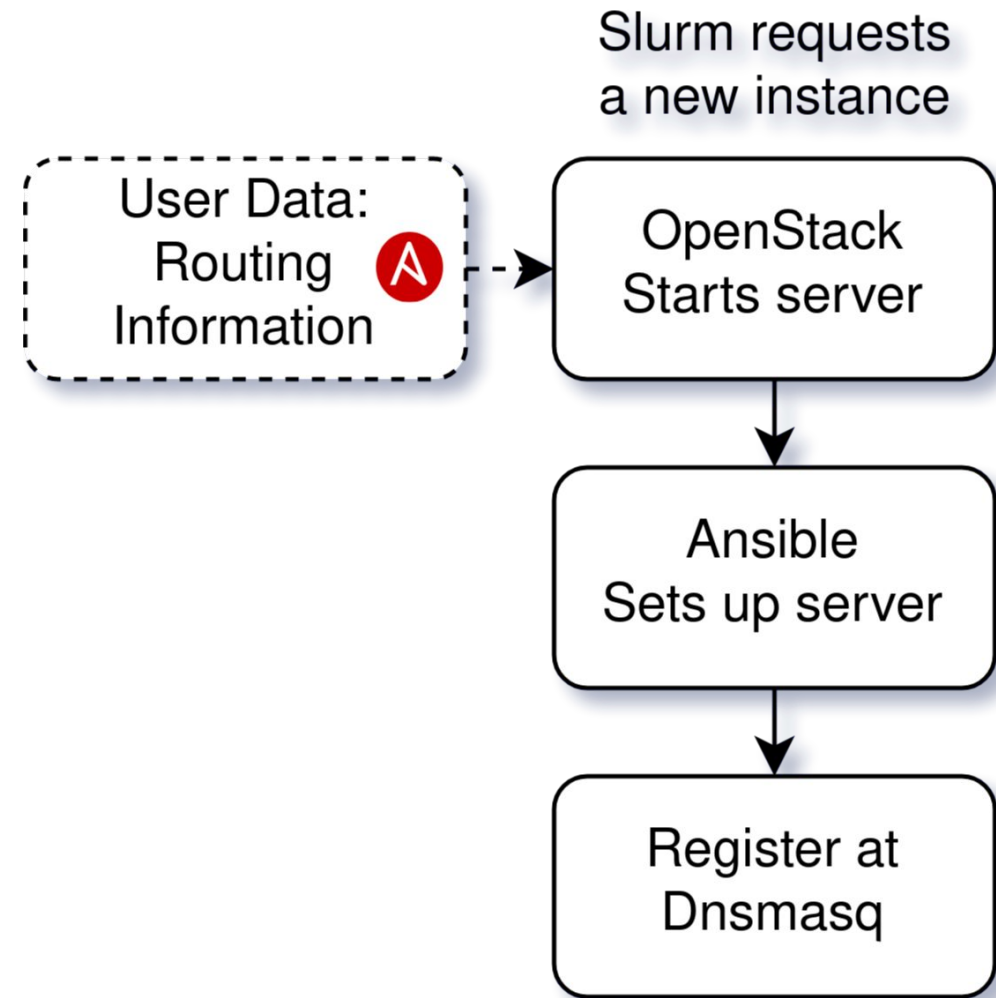
Workers from foreign clouds must “learn” how to communicate with the cluster’s master

=> **Add ip routes** over the respective vpn-gateway to reach the master

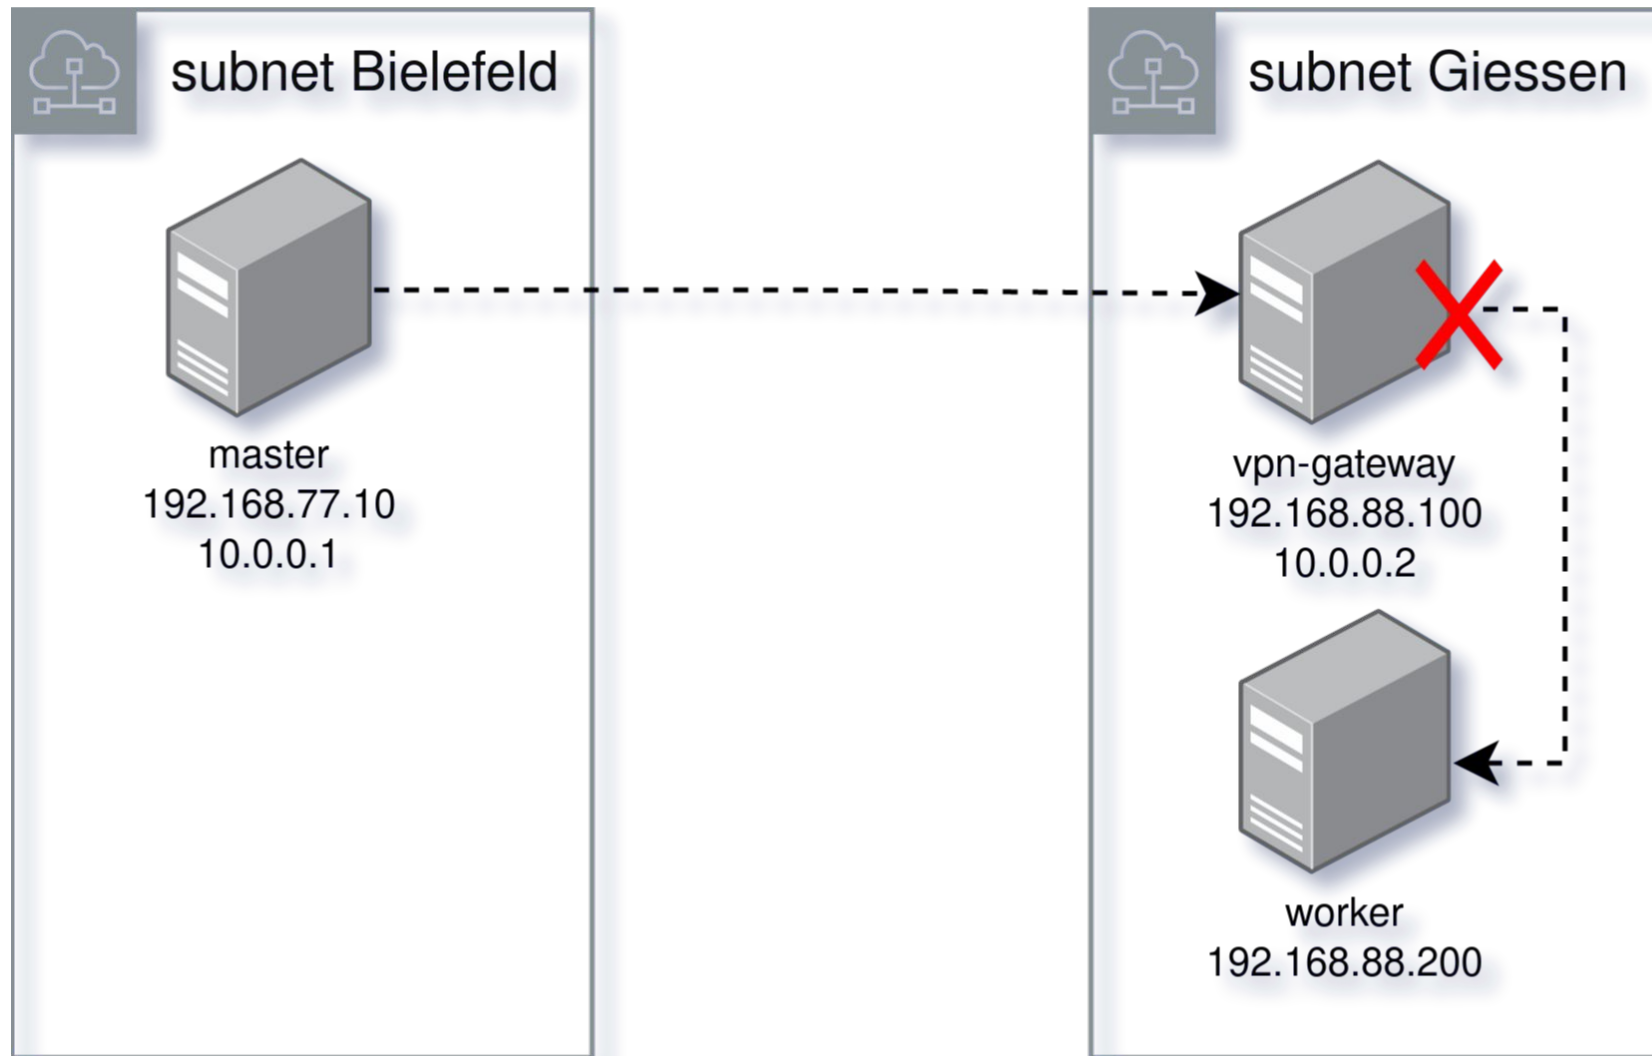
These routes must be injected **before Ansible** is executed as a route back must already exist at that point

=> Add ip routes using **User Data**

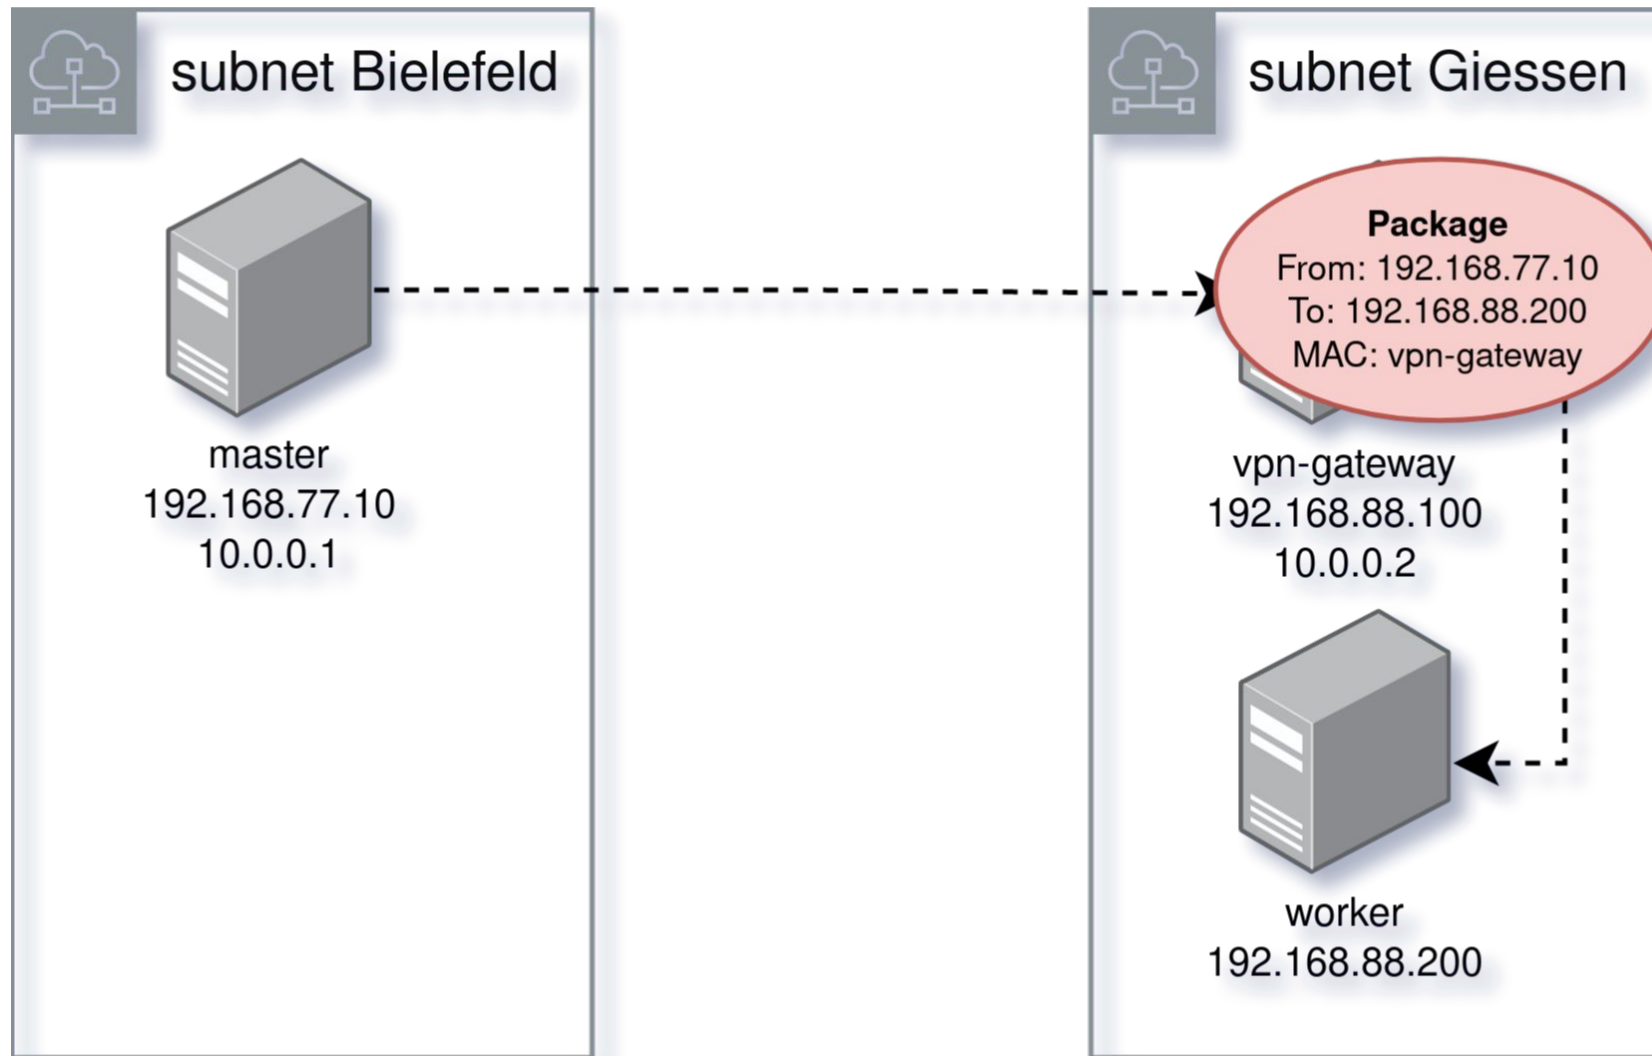
Vpn-gateways must **forward** ips forward over the default network



OpenStack: Port Security & VPN



OpenStack: Port Security & VPN



Multi-Cloud – BiBiGrid

- Cloud locations must be able to **communicate**
- Slurm must be able to request server instances for jobs “on demand” **on all clouds**
- Slurm must be able to schedule jobs to **specific clouds**
- Communication between clouds must be **secure**

- Go easy on resources
- Working on multiple clouds should not complicate use

Multi-Cloud – BiBiGrid

- Cloud locations must be able to **communicate – Routing (User Data), Port Security, Security Groups**
- Slurm must be able to request server instances for jobs “on demand” **on all clouds – VPN**
- Slurm must be able to schedule jobs to **specific clouds – Partitions**
- Communication between clouds must be **secure – Wireguard**

- Go easy on resources – Only one vpn-gateway for every cloud
- Working on multiple clouds should not complicate use – DNS, VPN

Outlook

- **Finalize** multi-cloud setup automation
- Test BiBiGrid's multi-cloud solution on an **actual use case** using **nextflow**
- Investigate **more efficient file sharing**

Outlook

- **Finalize** multi-cloud setup automation
- Test BiBiGrid's multi-cloud solution on an **actual use case**
- Investigate **more efficient file sharing**

=> Publish a BiBiGrid multi-cloud **Hands-on**



Multi-Cloud – BiBiGrid



[ELIXIR Germany](#)
Xaver Stiensmeier, Jan Krüger
www.elixir-europe.org

